



INDRAPRASTHA INSTITUTE OF INFORMATION TECHNOLOGY
NEW DELHI

Department of Computer Science & Engineering

CSE 343/ECE 363 : Machine Learning

Dr. Jainendra Shukla

Assignment - 4: CNN , K-Means

Anant Kumar Kaushal (2022067)

SECTION - A

A As per the question, input image is 'M x N' with P channels (dimensions)

(a) Dimensions of resulting feature map:-

$$\text{Output height} = \left\lfloor \frac{M - K + 2 \times P}{\text{stride}} \right\rfloor + 1$$

$$\text{Output width} = \left\lfloor \frac{N - K + 2 \times P}{\text{stride}} \right\rfloor + 1$$

Convolutional layer $\Rightarrow K \times K$

As stride = 1 & P = 0 :-

~~Feature map dimension = $[M - K + 1] [N - K + 1]$~~

\therefore Feature map dimension = (Output height) x (Output width)

$$= (M - K + 1) (N - K + 1)$$

(b) For each output pixel:

for P inputs, Total multiplications = $P \times K \times K = PK^2$

~~For~~ Total Addns = $P \times K \times K - 1 = PK^2 - 1$

Total = $2PK^2 - 1$

(c) As per (a), total number of output pixels per kernel = $(M - K + 1) (N - K + 1)$

For 'Q' kernels = $Q \times (M - K + 1) (N - K + 1)$

from (b), ^{order of} total no. of operations = $P \times K^2$
per pixel

Order of total ops. = $(P \times K^2) [O \times (M-K+1) \times (N-K+1)]$

= $O P K^2 (M-K+1) (N-K+1)$

$\therefore \min(M, N) \gg K :-$

$M-K+1 \approx M$

$N-K+1 \approx N$

$\therefore = O P K^2 M N$

Final complexity = $O(O \times P \times K^2 \times M \times N)$

Ⓐ B I Assignment step:

1. Distance calculation: For each point x_i , calculate distance b/w x_i and centroid.

$D(x_i, \mu_k) = |x_i - \mu_k|^2$

2. Cluster assignment: To the cluster, we assign new data point where μ_k is closest.

$C_k = \{x_i : |x_i - \mu_k|^2 \leq |x_i - \mu_j|^2 \forall j, 1 \leq j \leq k\}$

II Update Step:

For each cluster c_k , the new centroid μ_k as the mean of all data points assigned to that cluster.

$$\mu_k = \frac{1}{|C_k|} \sum_{x_i \in C_k} x_i$$

The 'Elbow method' is a commonly used technique in clustering analysis to determine the optimal number of clusters (k) for a dataset when using algorithms like k -means.

Within-Cluster Sum of Squares (WCSS):

The WCSS measures the sum of squared distances between each data point and the centroid of its cluster. It quantifies how compact the clusters are.

$$WCSS = \sum_{k=1}^K \sum_{x \in C_k} \|x - \mu_k\|^2$$

K : No's of clusters

C_k : k^{th} Cluster

μ_k : Centroid of k^{th} cluster

$\|x - \mu_k\|^2$: Squared Euclidean distance

After the above computation, the plot would show a decreasing graph with increasing k (on the x -axis) and WCSS on the y -axis).

The new point can be identified as one where the rate of decrease in WCSS sharply changes.

Yes, random initialization can sometimes arrive at the global minimum, especially if the initial centroids happen to be well-placed. Randomly assigning cluster leads to convergence but it is not guaranteed to find the global minimum. Different initial centroids can lead to different clustering results hence converging to local minimum (not the global one).

$$WCSS = \sum_{i=1}^k \sum_{x \in C_i} ||x - \mu_i||^2$$

For each cluster C_i , the centroid μ_i is calculated as the mean of the points in the cluster.

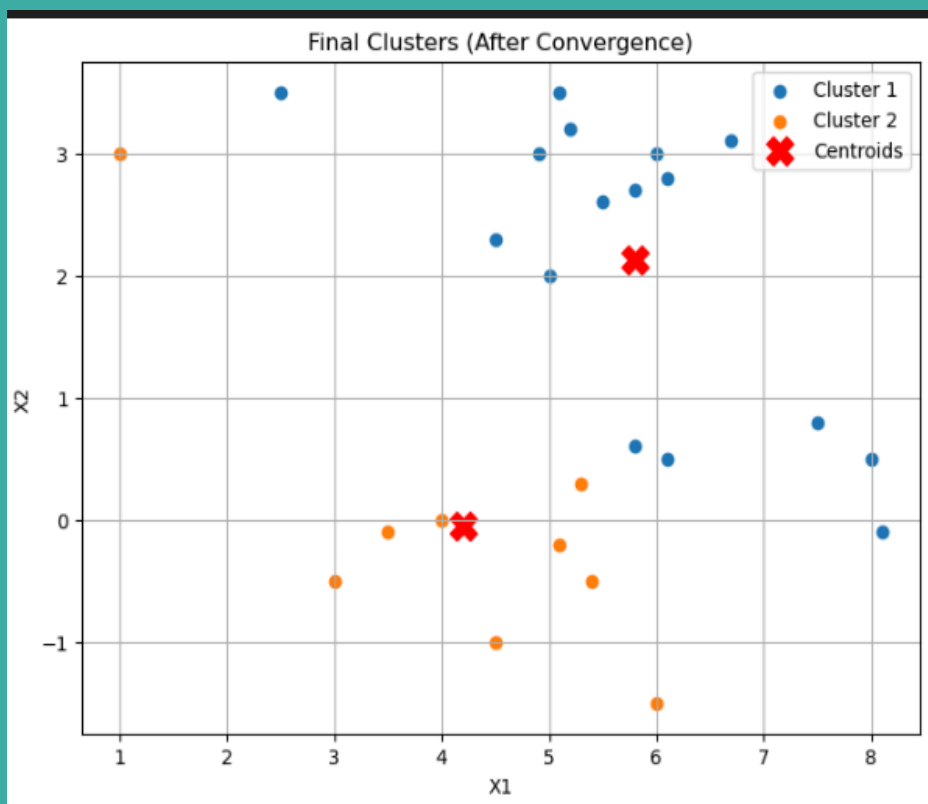
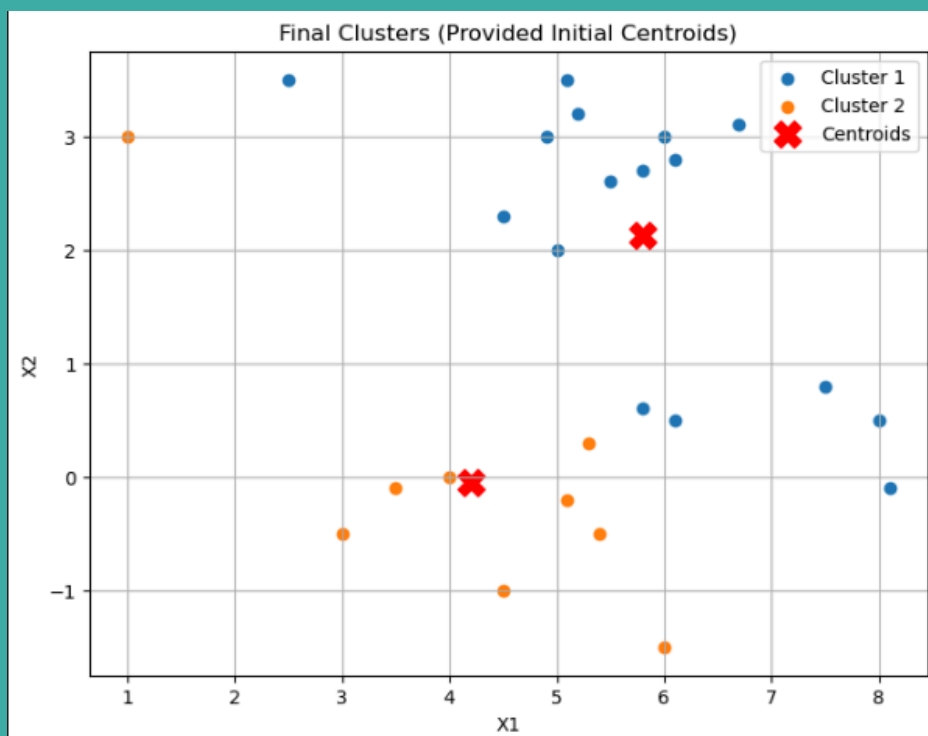
SECTION - B

1. Implementation of K-Means Clustering from scratch including all important parameters is as follows:-

```
def kmeans_clustering(X, centroids, max_iter=100, threshold=1e-4):  
    for iteration in range(max_iter):  
        distances = np.linalg.norm(X[:, np.newaxis] - centroids, axis=2)  
        labels = np.argmin(distances, axis=1)  
        new_centroids = np.array([X[labels == k].mean(axis=0) if len(X[labels  
== k]) > 0 else centroids[k] for k in range(len(centroids))])  
  
        if np.linalg.norm(new_centroids - centroids) < threshold:  
            print(f"Converged after {iteration + 1} iterations.")  
            break  
        centroids = new_centroids  
    return centroids, labels
```

2. Final Values after algorithm convergence and the useful cluster plots are as follows:-

```
Final Centroids:  
[[ 5.8         2.125        ]  
 [ 4.2        -0.05555556]]
```

3. Comparison of the results using the provided initial centroids versus using random initialization of centroids is as follows :-

Both approaches produced the same number of final clusters, and both converged at the same number of iterations. This suggests that the two initial centroid selection methods are not very different in the impact on the final clusters obtained in this particular data set.



But again, it has to be noted that such might not always be the case. For other data sets, another initialization method could mean that different clustering results can be obtained, possibly necessitating more iterations to converge.

```
Final Centroids (Provided Initial Centroids):
```

```
[[ 5.8          2.125         ]  
 [ 4.2          -0.05555556]]
```

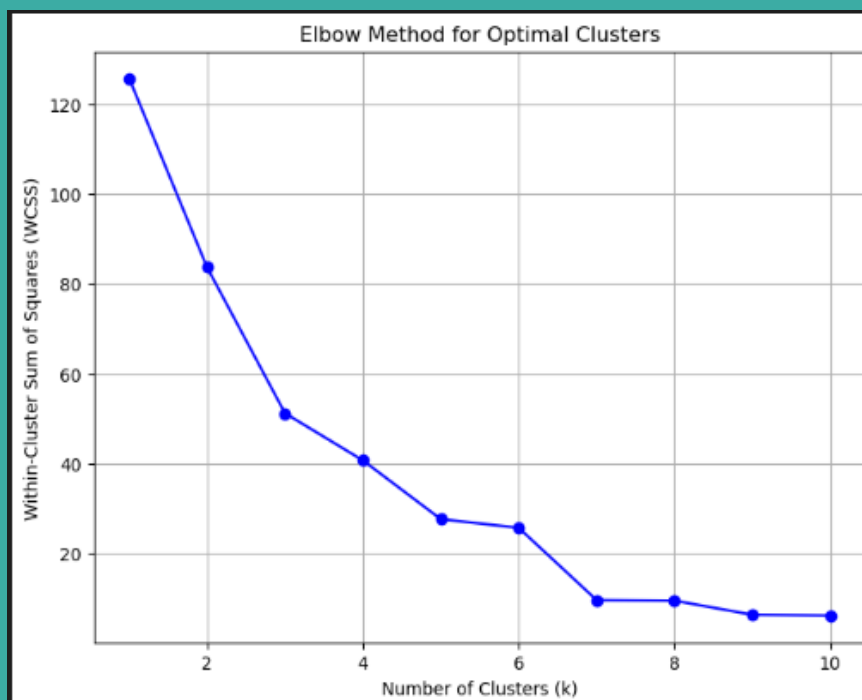
```
Final Centroids (Random Initial Centroids):
```

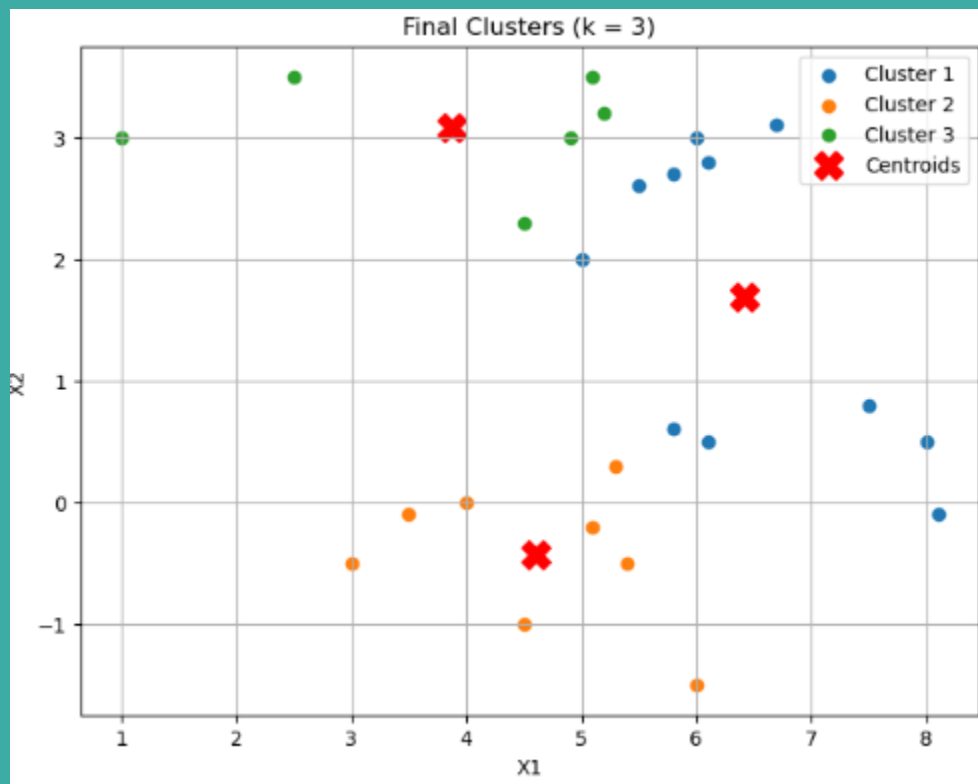
```
[[ 5.8          2.125         ]  
 [ 4.2          -0.05555556]]
```

4. Elbow Method , WCSS & Final Cluster Plot :-

As we know, the elbow method is a method to figure out the ideal number of clusters in a data set when doing k-means clustering. It involves plotting the within-cluster sum of squares, or WCSS, versus the number of clusters. The **elbow point** is the point where the curve bends, giving the best number of clusters to minimize variance without overfitting. So if we see the WCSS curve for this question we see the elbow is made at **k = 3** therefore it is the optimal value of k . After this optimal value of k, there is no significant decrease in the y-axis values compared to before.





SECTION - C

1. Data Preparation:-

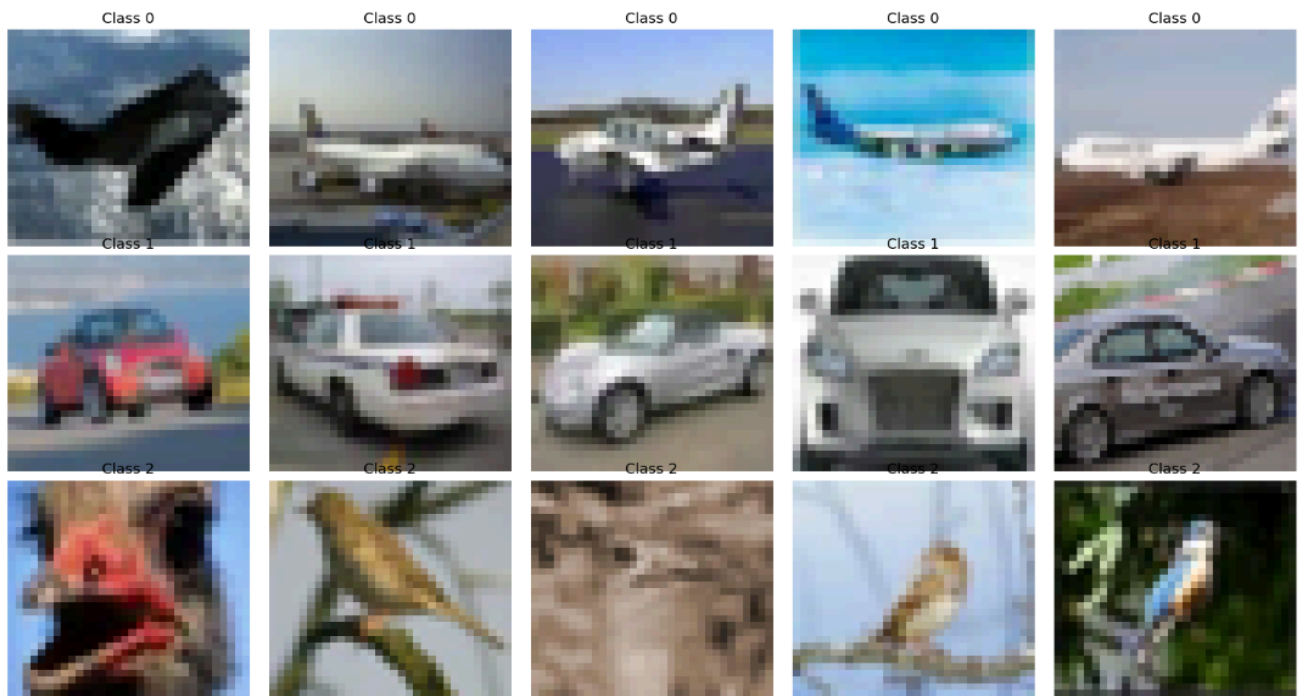
I created a custom Dataset class for the data and created data loaders for all the dataset splits - train, val, and test. Here, the 15,000 images from the training dataset are split into train-val via 80:20 split, and 3,000 images are retained as the testing data from the original test dataset . The output for the same is as follows:-

```
Train size: 12000  
Validation size: 3000  
Test size: 3000
```


2. Visualization:-

The image of each class for both training and validation data can be found below:-

Training Data: 5 Images per Class



Validation Data: 5 Images per Class



3. CNN Implementation:-

I have implemented the following CNN class satisfying all requirements with proper flattening that was asked in the question:-

```
class CustomCNN(nn.Module):
    def __init__(self, num_classes=3):
        super(CustomCNN, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=16, kernel_size=5, stride=1, padding=1)
        self.pool1 = nn.MaxPool2d(kernel_size=3, stride=2)
        self.conv2 = nn.Conv2d(in_channels=16, out_channels=32, kernel_size=3, stride=1, padding=0)
        self.pool2 = nn.MaxPool2d(kernel_size=3, stride=3)
        self.fc1 = nn.Linear(32 * 4 * 4, 16)
        self.fc2 = nn.Linear(16, num_classes)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = self.pool1(x)
        x = F.relu(self.conv2(x))
        x = self.pool2(x)
        x = torch.flatten(x, 1)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x

num_classes = len(chosen_classes)
model = CustomCNN(num_classes=num_classes)
print(model)
```

✓ 0.0s

```
CustomCNN(
  (conv1): Conv2d(3, 16, kernel_size=(5, 5), stride=(1, 1), padding=(1, 1))
  (pool1): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
  (pool2): MaxPool2d(kernel_size=3, stride=3, padding=0, dilation=1, ceil_mode=False)
  (fc1): Linear(in_features=512, out_features=16, bias=True)
  (fc2): Linear(in_features=16, out_features=3, bias=True)
)
```

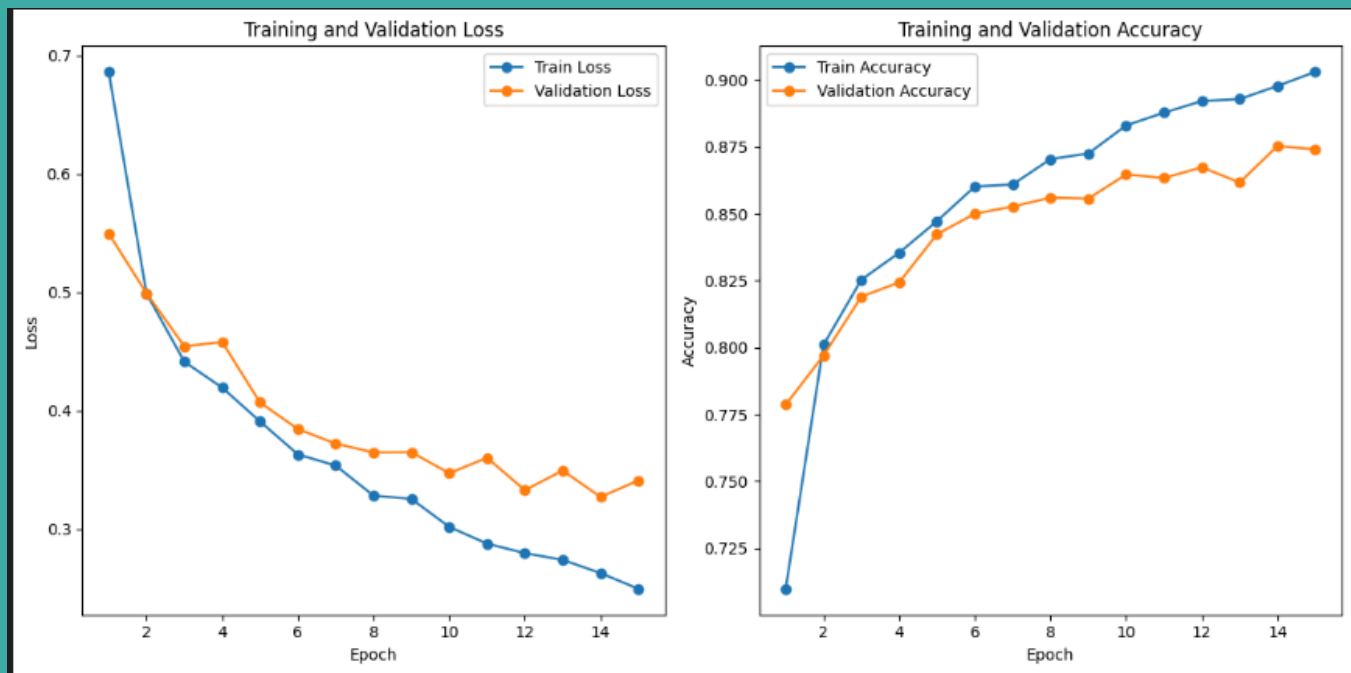
4. Model Training(CNN) :-

Implementation of the above CNN model has been done for all 15 epochs and all performance metrics has been recorded which can be found below :-

```
Model already trained and loaded.
Epoch [1/15] - Train Loss: 0.6863, Train Accuracy: 0.7098 - Val Loss: 0.5498, Val Accuracy: 0.7787
Epoch [2/15] - Train Loss: 0.4991, Train Accuracy: 0.8010 - Val Loss: 0.4993, Val Accuracy: 0.7970
Epoch [3/15] - Train Loss: 0.4414, Train Accuracy: 0.8253 - Val Loss: 0.4545, Val Accuracy: 0.8190
Epoch [4/15] - Train Loss: 0.4196, Train Accuracy: 0.8354 - Val Loss: 0.4581, Val Accuracy: 0.8243
Epoch [5/15] - Train Loss: 0.3910, Train Accuracy: 0.8472 - Val Loss: 0.4071, Val Accuracy: 0.8423
Epoch [6/15] - Train Loss: 0.3632, Train Accuracy: 0.8602 - Val Loss: 0.3846, Val Accuracy: 0.8500
Epoch [7/15] - Train Loss: 0.3537, Train Accuracy: 0.8609 - Val Loss: 0.3723, Val Accuracy: 0.8527
Epoch [8/15] - Train Loss: 0.3283, Train Accuracy: 0.8704 - Val Loss: 0.3649, Val Accuracy: 0.8560
Epoch [9/15] - Train Loss: 0.3256, Train Accuracy: 0.8725 - Val Loss: 0.3651, Val Accuracy: 0.8557
Epoch [10/15] - Train Loss: 0.3019, Train Accuracy: 0.8830 - Val Loss: 0.3473, Val Accuracy: 0.8647
Epoch [11/15] - Train Loss: 0.2878, Train Accuracy: 0.8878 - Val Loss: 0.3603, Val Accuracy: 0.8633
Epoch [12/15] - Train Loss: 0.2797, Train Accuracy: 0.8921 - Val Loss: 0.3329, Val Accuracy: 0.8673
Epoch [13/15] - Train Loss: 0.2741, Train Accuracy: 0.8928 - Val Loss: 0.3495, Val Accuracy: 0.8617
Epoch [14/15] - Train Loss: 0.2629, Train Accuracy: 0.8978 - Val Loss: 0.3272, Val Accuracy: 0.8753
Epoch [15/15] - Train Loss: 0.2496, Train Accuracy: 0.9030 - Val Loss: 0.3412, Val Accuracy: 0.8740
```


5. Observation (CNN) :-



Confusion matrix , plots , F1-score and accuracy can be found below :-



Test Accuracy: 0.8860

Test F1-Score: 0.8855





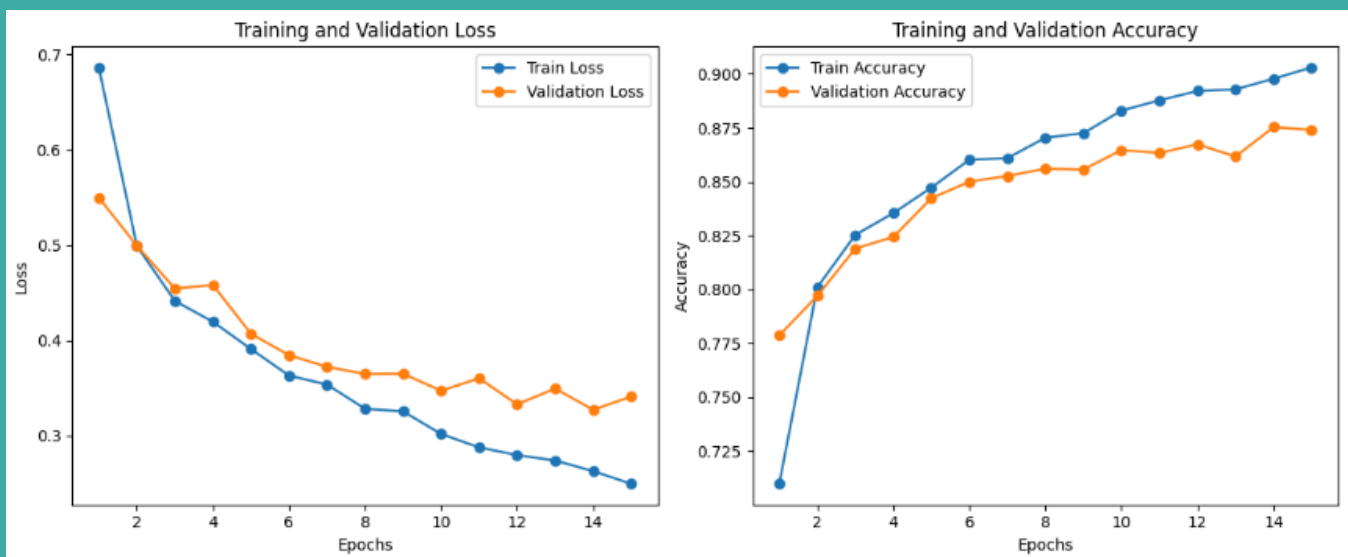
6. Model Training(MLP) :-

Implementation of the above MLP model has been done for all 15 epochs and all performance metrics has been recorded which can be found below :-

```
Model already trained and loaded.
Training and Validation Metrics:
Epoch [1/15] - Train Loss: 0.7128, Train Accuracy: 0.6973 - Val Loss: 0.5457, Val Accuracy: 0.7777
Epoch [2/15] - Train Loss: 0.5058, Train Accuracy: 0.7965 - Val Loss: 0.4800, Val Accuracy: 0.8087
Epoch [3/15] - Train Loss: 0.4494, Train Accuracy: 0.8213 - Val Loss: 0.4479, Val Accuracy: 0.8227
Epoch [4/15] - Train Loss: 0.4112, Train Accuracy: 0.8408 - Val Loss: 0.4315, Val Accuracy: 0.8273
Epoch [5/15] - Train Loss: 0.3857, Train Accuracy: 0.8456 - Val Loss: 0.4321, Val Accuracy: 0.8360
Epoch [6/15] - Train Loss: 0.3667, Train Accuracy: 0.8556 - Val Loss: 0.3843, Val Accuracy: 0.8473
Epoch [7/15] - Train Loss: 0.3475, Train Accuracy: 0.8623 - Val Loss: 0.3684, Val Accuracy: 0.8567
Epoch [8/15] - Train Loss: 0.3283, Train Accuracy: 0.8703 - Val Loss: 0.3847, Val Accuracy: 0.8503
Epoch [9/15] - Train Loss: 0.3227, Train Accuracy: 0.8715 - Val Loss: 0.3528, Val Accuracy: 0.8687
Epoch [10/15] - Train Loss: 0.3085, Train Accuracy: 0.8794 - Val Loss: 0.3365, Val Accuracy: 0.8693
Epoch [11/15] - Train Loss: 0.2971, Train Accuracy: 0.8842 - Val Loss: 0.3408, Val Accuracy: 0.8653
Epoch [12/15] - Train Loss: 0.2908, Train Accuracy: 0.8858 - Val Loss: 0.3363, Val Accuracy: 0.8733
Epoch [13/15] - Train Loss: 0.2802, Train Accuracy: 0.8918 - Val Loss: 0.3412, Val Accuracy: 0.8687
Epoch [14/15] - Train Loss: 0.2703, Train Accuracy: 0.8942 - Val Loss: 0.3188, Val Accuracy: 0.8783
Epoch [15/15] - Train Loss: 0.2593, Train Accuracy: 0.9000 - Val Loss: 0.3676, Val Accuracy: 0.8603
```


7. Observation (MLP) :-



Confusion matrix , plots , F1-score and accuracy can be found below :-



Test Accuracy: 0.7967

Test F1-Score: 0.7964





8. Comparison :-

Since the plots are already posted above so the key insights of them is written below :-

The MLP model has a steady improvement in accuracy and decreasing loss over epochs, which means it is learning effectively. However, the validation loss is still higher than the training loss, which indicates some overfitting. The CNN model has a higher accuracy than the MLP model, which means that it generalizes better. The training and validation loss curves for the CNN model are **more stable and consistent**, with less overfitting compared to the MLP model.

```
MLP Model - Test Accuracy: 0.7967, F1-score: 0.7964
```

```
CNN Model - Test Accuracy: 0.8860, F1-score: 0.8855
```

